

Robot Perception and Control


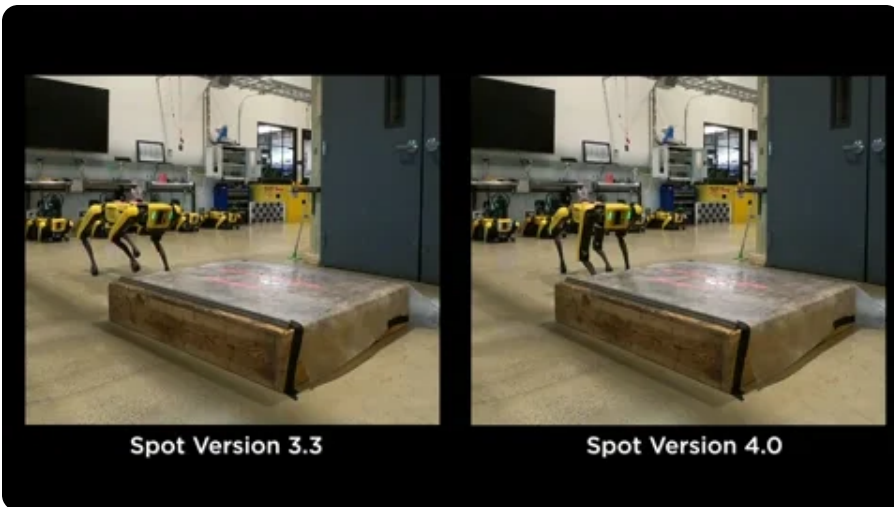
Legged Locomotion

Last updated: Jul / 25 /2024

Kashu Yamazaki

kyamazak@andrew.cmu.edu

What can we do with RL × Legged Robots?



Actuator Networks arxiv ↗

Actuators are extremely difficult to model accurately.

- *nonlinear* and *non-smooth* dissipation in dynamics.
- contains cascaded feedback loops and a number of internal states that are not even directly observable.

Actuator Networks is a **data driven** solution that can provide better simulation of an actuator via supervised learning.

- learns **action-to-torque** relationship that includes all software and hardware dynamics.
- actuator network estimated torque at the joints given a history of position errors and velocities.




Fig. 1. Creating a control policy. In the first step, we identify the physical parameters of the robot and estimate uncertainties in the identification. In the second step, we train an actuator net that models complex actuator/software dynamics. In the third step, we train a control policy using the models produced in the first two steps. In the fourth step, we deploy the trained policy directly on the physical system.

“ collect *joint position errors*, *velocities*, and *torque* using a controller for more than a million samples with varied amplitude and frequency and manual disturbances for diverse situation.

Learning by Cheating [arxiv ↗](#) [github ↗](#)

Proposed two-stage training procedure, which first train a privileged agent and then using the agent as a teacher to train a purely vision-based system, for effective imitation learning. This paradigm is the underlying concept in the legged RL.

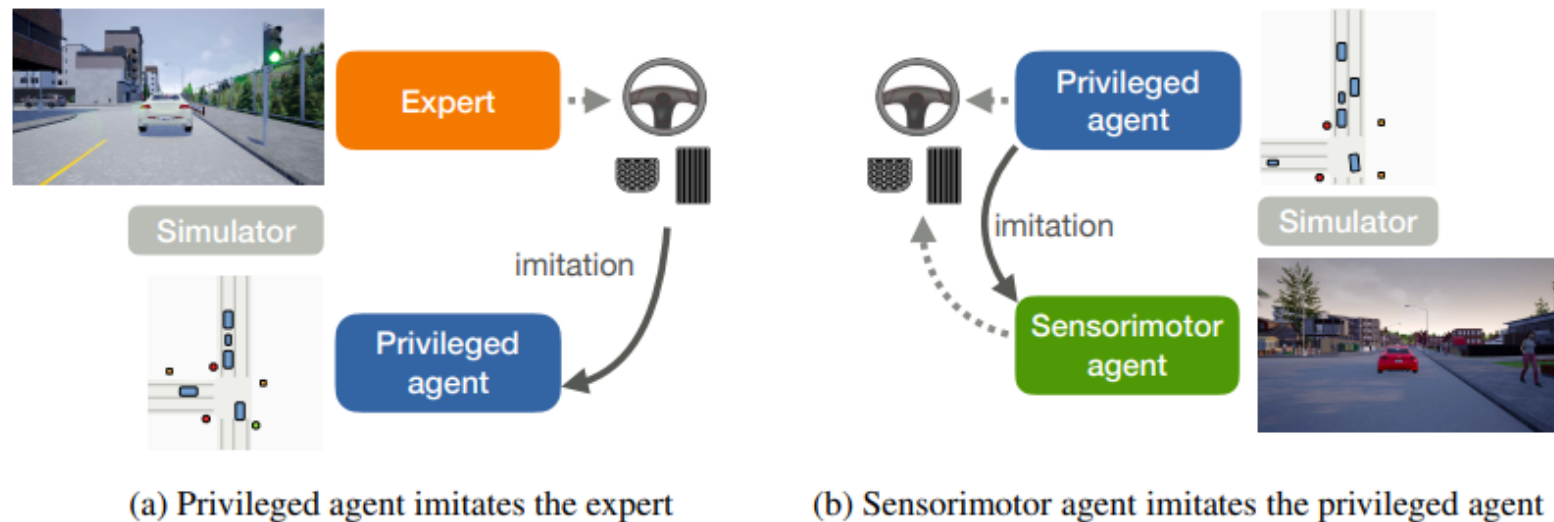
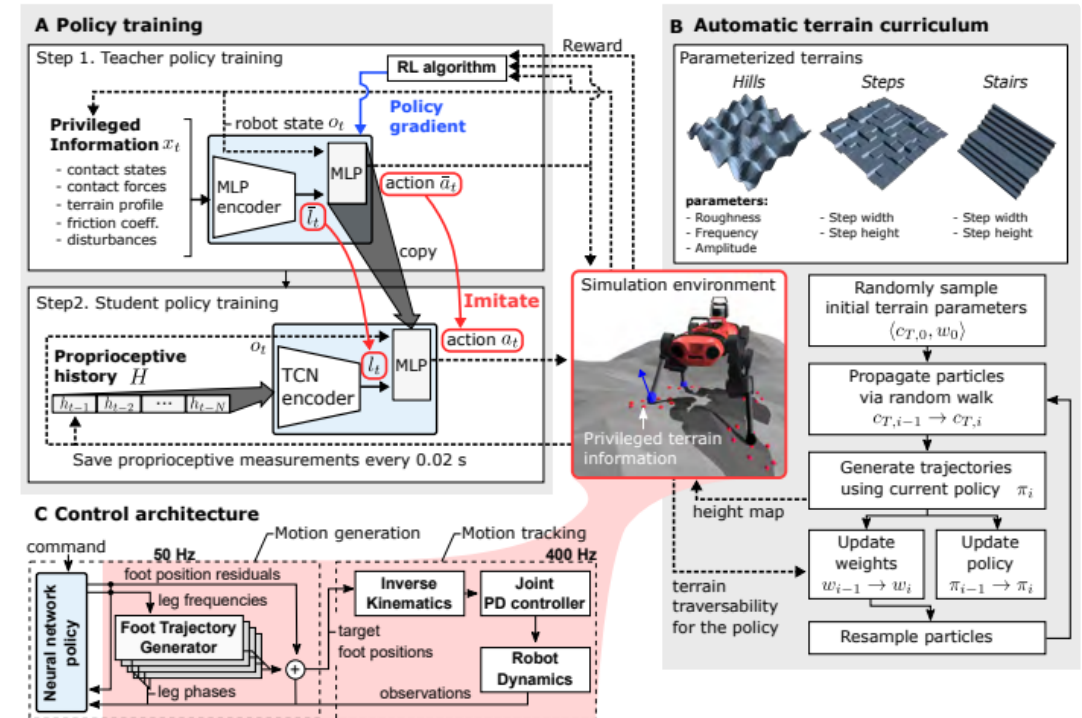


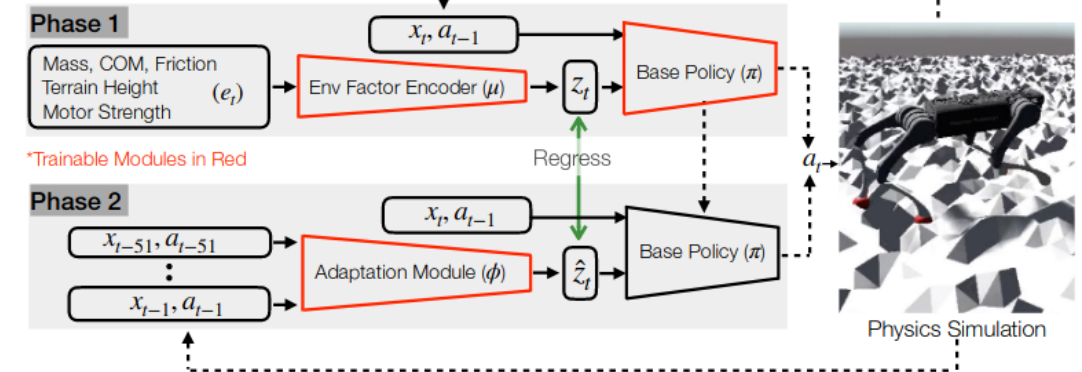
Figure 1: Overview of our approach. **(a)** An agent with access to privileged information learns to imitate expert demonstrations. This agent learns a robust policy by cheating. It does not need to learn to see because it gets direct access to the environment’s state. **(b)** A sensorimotor agent without access to privileged information then learns to imitate the privileged agent. The privileged agent is a “white box” and can provide high-capacity on-policy supervision. The resulting sensorimotor agent does not cheat.

Learning Locomotion over Challenging Terrain [arxiv ↗](#) [github ↗](#)




RMA: Rapid Motor Adaptation [paper ↑](#)

A) Training in Simulation




B) Deployment



Learning to Walk in Minutes [arxiv ↗](#) [github ↗](#)

Presents a training setup that achieves fast policy generation for real-world robotic tasks by using massive parallelism on a single workstation GPU (showcase of Isaac Gym).

- A codebase is widely used as baseline for developing legged locomotion system.



Walk These Ways arxiv ↗

Perceptive locomotion


Perceptive locomotion for quadrupeds

Presented a three stage training and deploy method to perform zero-shot sim-to-real transfer [1 ↗].

1. a **teacher policy**, which has access to privileged information, is trained to follow a random target velocity over randomly generated terrain with random disturbances.
2. a **student policy** is trained to reproduce the teacher policy's actions without using this privileged information.
3. transfer the learned student policy to the physical robot and deploy it in the real world with onboard sensors.

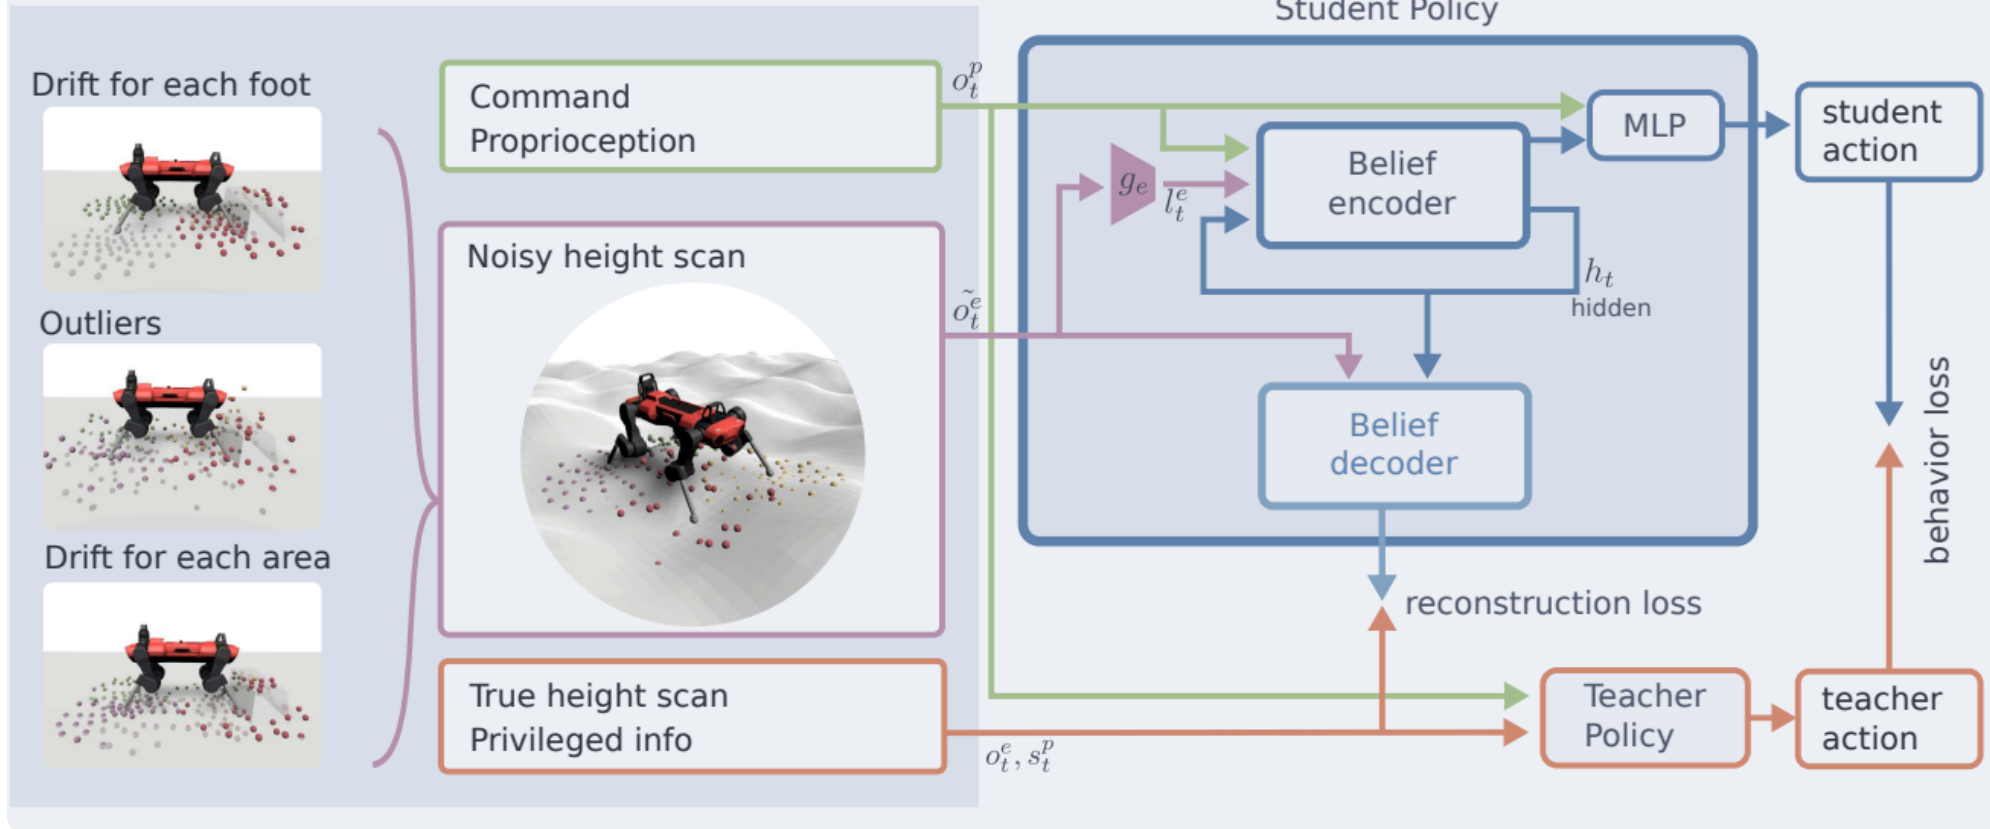
Training teacher policy

1. Teacher policy training

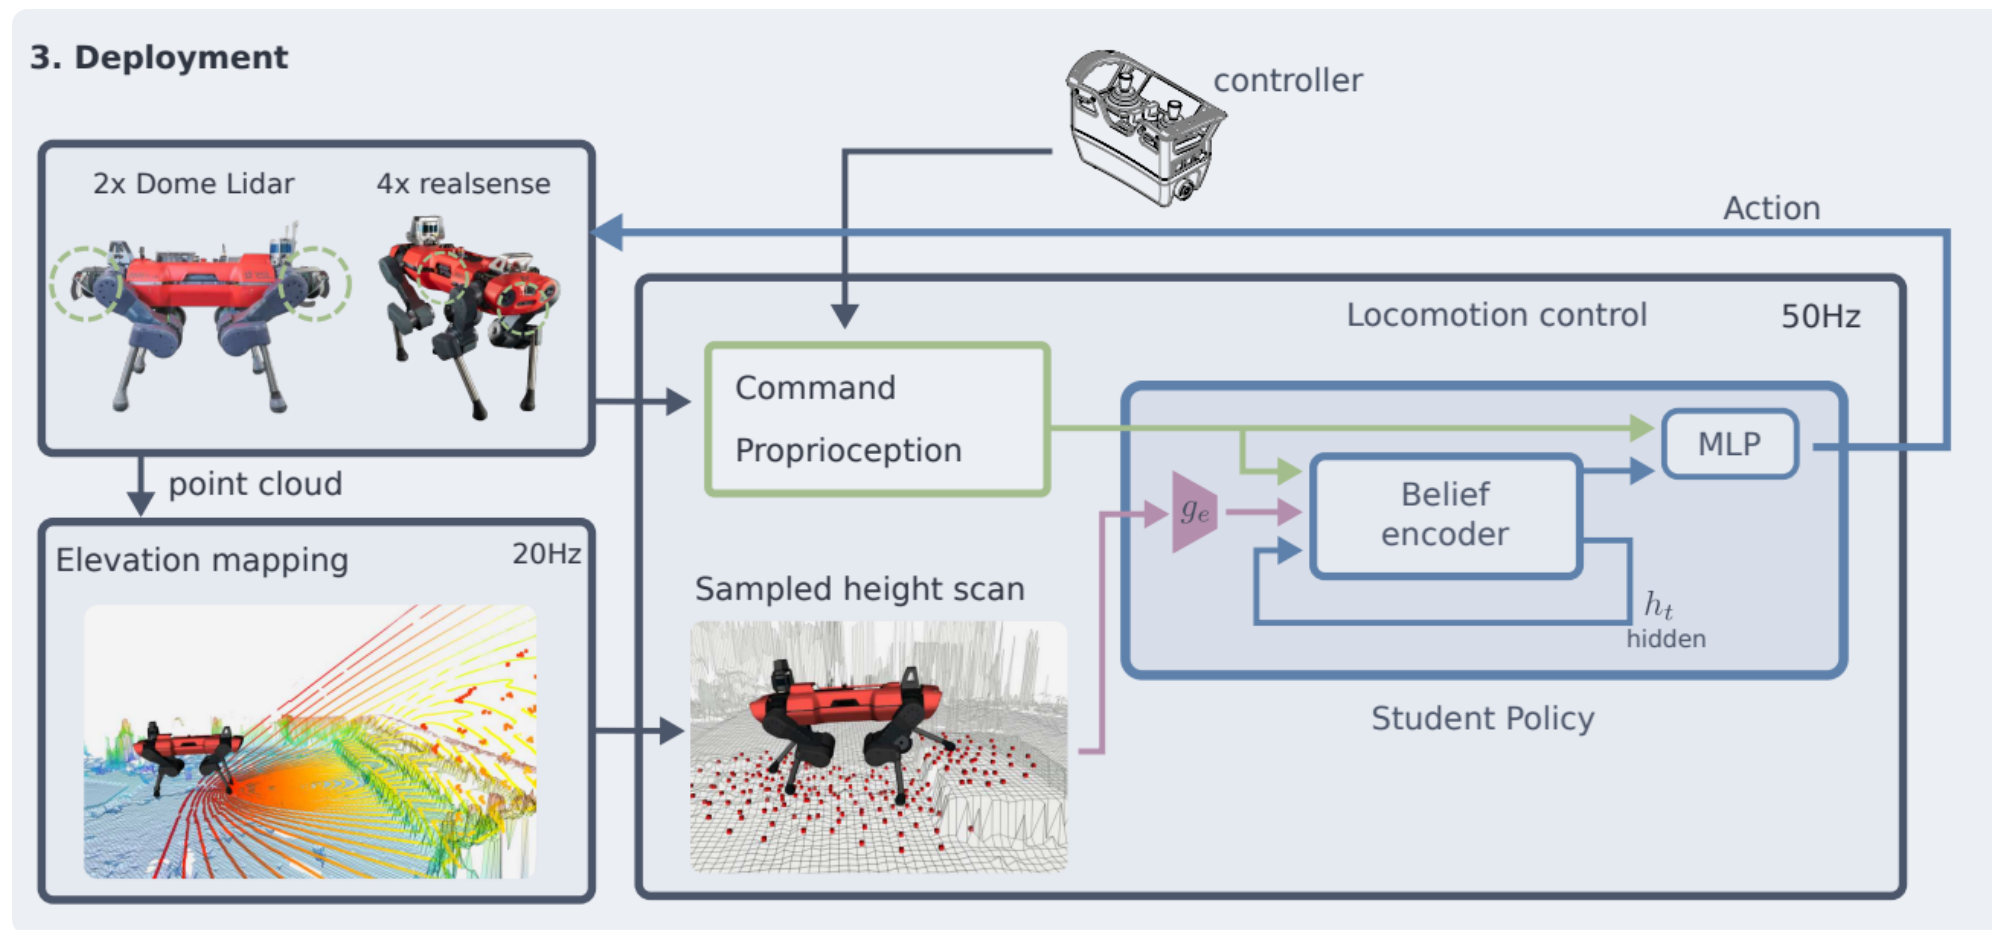


Training student policy

2. Student policy training



Deployment



Legged Locomotion using Egocentric Vision [arxiv ↗](#)

Upstairs
17cm high, 30cm deep



	Success	#Stairs
Ours	100%	13
Blind	0%	2.2

Downstairs
17cm high, 30cm deep



	Success	#Stairs
Ours	100%	13
Blind	100%	13

Stepping Stones
30cm wide, 15cm apart



	Success	#Stones
Ours	94%	9.4
Blind	0%	0


Gaps
26cm apart



	Success
Ours	100%
Blind	0%

Parkour Learning arxiv ↗

Extreme Parkour arxiv ↗



Humanoid Parkour Learning arxiv ↗